
The background of the slide features a yellow Pikachu on the right side, looking surprised with its mouth open. It is surrounded by several blue, shadowed shapes of different Pokémons, including Eevees and Pichus. Red lightning bolts are scattered across the background. Large question marks are placed at the top and bottom corners of the slide.

Comparative Analysis of CNN Architectures on Pokémon Image Classification



Project Overview & Goals

- ❖ Evaluate 3 convolutional neural network (CNN) architectures on a Pokémon image dataset
 - SimpleCNN, VGG11, and ResNet18
- ❖ Train models with varied hyperparameters to assess performance, convergence, and generalization.
- ❖ Goals:
 - Compare architectural impacts on learning dynamics and robustness.
 - Analyze hyperparameter effects on optimization and stability.
 - Derive insights into transferable design choices for image classification tasks.

Dataset Overview

- ❖ Dataset: Pokémon images split into train, validation, and test sets.
 - Batch sizes: 128 train, 256 val, 256 test
- ❖ 110 classes representing distinct Pokémon species (e.g., Porygon, Goldeen, Hitmonlee, ..., Raichu).
- ❖ Image specifications: 64×64 resolution, RGB format



- ❖ Preprocessing:
 - Training: random resized crop, horizontal flip, normalization.
 - Validation/Test: resize, center crop, normalization.
 - Rationale: Diverse visual features (shapes, colors, poses) support fine-grained classification analysis.

Architecture 1: SimpleCNN

- ❖ Lightweight baseline CNN with stacked 3×3 convolutional layers, batch normalization, ReLU activations, and max pooling.
 - 3 conv blocks ($64 \rightarrow 128 \rightarrow 256$), each with Conv–BN–ReLU–Conv–BN–ReLU–MaxPool
- ❖ Fully connected classifier with dropout, configured for 110-class output.
 - 3 conv blocks ($64 \rightarrow 128 \rightarrow 256$), each with Conv–BN–ReLU–Conv–BN–ReLU–MaxPool
- ❖ Focuses on basic feature extraction without advanced mechanisms.
- ❖ Why we chose this: Serves as a control to highlight limitations of shallow networks on complex Pokémon images.




Architecture 2: VGG11

- ❖ Deeper network using uniform 3×3 convolutions, batch normalization, ReLU, and max pooling.
 - Two fully connected layers with dropout, ending in a 110-class classifier.
 - Dense head ($512 \times 4 \times 4 \rightarrow 512 \rightarrow 512 \rightarrow 110$ with dropout)
- ❖ Emphasizes depth through repetitive small-kernel stacking.
- ❖ Why we chose this: Demonstrates benefits and challenges of deeper plain CNNs (e.g., vanishing gradients).



Architecture 3: ResNet18

- ❖ Residual network with 18 layers and shortcut connections for identity mappings.
 - 4 stages (64, 128, 256, 512) + global average pooling
→ 110-way final linear layer
 - Improved gradient flow via residual blocks with convolutions and batch normalization.
- ❖ Why we chose this: Illustrates advantages of residual learning for deeper training and better generalization.




Experimental Setup

- ❖ Models: SimpleCNN, VGG11, ResNet18
- ❖ Hyperparameter configurations:

Config Name	Optimizer	LR	Momentum	Weight Decay	Scheduler	Epochs
SGD_StepDecay	SGD	0.1	0.9	5e-4	StepLR	200
SGD_Cosine	SGD	0.1	0.9	5e-4	CosineAnnealingLR	200
Adam_NoReg	Adam	0.001	N/A	0.0	None	100
Adam_WD	Adam	0.0005	N/A	1e-4	ReduceLROnPlateau	150


Learning Curves SimpleCNN

- ❖ Training loss steadily decreases
- ❖ Validation loss consistently increases
- ❖ Validation accuracy remains low and highly unstable
- ❖ SGD (StepDecay / CosineAnnealing) shows slightly more stable curves but no significant improvement in accuracy
- ❖ Adam yields faster training loss reduction but no generalization gains



Simple CNN - Confusion Matrix


- ❖ Some configurations (SGD_StepDecay, SGD_Cosine) exhibit mode collapse, repeatedly predicting only a handful of classes
- ❖ Adam_NoReg and Adam_WD scatter predictions across many wrong categories, essentially uniform guessing.



Learning Curves


VGG11

- ❖ All VGG11 learning curves show:
 - Train loss dropping steadily
 - Validation loss rising over time
 - Validation accuracy fluctuating near random
 - This pattern shows the model memorizes training images but fails to learn class-discriminative features that transfer to validation images.



VGG11 - Confusion Matrix


- ❖ Consistently fails to generalize on this dataset
 - Across all hyperparameter configurations validation accuracy stays between 0–6%
 - Most predictions fall into only 4–10 output classes (out of 110).
 - Adam configs scatter predictions across many incorrect classes (uniform guessing)



Learning Curves


ResNet18

- ❖ Training loss decreases smoothly across all optimizers, indicating successful optimization and effective use of residual connections.
 - Validation accuracy is unstable and low (0.5–2.8%), but consistently higher than SimpleCNN and VGG11, meaning ResNet18 extracts more useful features.
 - SGD produces the highest test accuracies (1–2.2%) but still fluctuates heavily.
 - Adam configs train fastest but generalize poorly, especially without weight decay.
- ❖ ResNet18 learns more robust representations than the other two models, but still struggles with the 110-class Pokémon dataset.




ResNet18 Confusion Matrix

- ❖ Many predictions concentrate in a band across certain classes, but the model fails to consistently match true labels.
- ❖ Compared to SimpleCNN/VGG11, ResNet18 still cannot separate 110 visually similar classes on this limited dataset.



Key Findings



❖ All three models struggled significantly on this fine-grained 110-class Pokémon dataset when trained from scratch.

- ❖ Optimizer & scheduler choice mattered, with Adam variants producing higher score for SimpleCNN and VGG11, and SGD variants producing highest test scores for ResNet18.
- ❖ Confusion matrices lacked diagonal structure, indicating that models weren't learning meaningful class-discriminative features.



Cross-Dataset Insights

- ❖ Architectures that work well on ImageNet do not transfer well when trained from scratch on small datasets.
- ❖ Hyperparameters that succeed on CIFAR/ImageNet (SGD + cosine decay) did not translate into meaningful gains here.
- ❖ Results imply that transfer learning or pretraining is essential for small-scale, high-class image tasks.

What Surprised Me

- ❖ Even deep models failed completely. We expected ResNet18 to achieve at least 20-30% accuracy, but it never learned meaningful decision boundaries.
- ❖ SGD sometimes performed better than Adam, despite slower initial convergence, showing that “fast loss decrease” doesn’t mean better generalization.
- ❖ Confusion matrices were essentially random, with no diagonal structure
- ❖ The dataset’s diversity + class imbalance had a much larger effect than anticipated.

Limitations & Future Work

❖ Limitations:

- Dataset small relative to number of classes; some classes had very few samples.
 - Training set had 110 classes, but the test set had (unannounced) 150 classes.
- Hyperparameter systematic tuning of LR, WD, and batch size.



❖ Future Work:

- Apply pretrained backbones and fine-tune; likely dramatic improvement.
- Try class-balanced sampling to stabilize training.

Thanks for listening!

- ❖ Any Questions?
- ❖ Pokemon dataset →
<https://huggingface.co/datasets/keremberke/pokemon-classification>

